

Variations on variants

Roi Barkan

Slides online [here](#)

Hi, I'm Roi

- Roi Barkan - רועי ברקן
- I live in Tel Aviv, Israel
- C++ developer since 2000
- VP Technologies @ Istra Research
 - Finance, Low Latency, in Israel
 - careers@istraresearch.com
- First time to CPPNow
 - Try to be active in chat, Q&A
 - Please - ask questions, make comments

Agenda

- Introduction
 - What is `std::variant`, what is it good for
 - Typical uses for `Variant`
- Variants vs. Unions
 - Examples
 - Existing Approaches
 - Pattern Matching
 - Variants with Commonality
- Intrusive Variants
- Streams of Variants
- Variants for Devirtualization

Introduction

What is a variant

- Cppreference.com:
 - The class template `std::variant` represents a type-safe **union**.
- Boost.org:
 - The variant class template is a safe, generic, stack-based *discriminated* **union** container.
- Plain English - a **union** that knows (holds) its type.

```
union MyUnion {  
    int integer;  
    double real;  
};  
//...  
void foo(const MyUnion& uni) {  
    cout << uni.integer << endl;  
} i.barkan@gmail.com
```

```
using MyVariant = variant<int,double>;  
//....  
void bar(const MyVariant& var) {  
    cout << std::get<int>(var);  
    //or  
    std::visit(  
        [](auto &item) {cout << item;}, var);  
}
```

Memory Layout

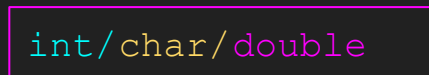
```
struct {int i; char c; double d};
```



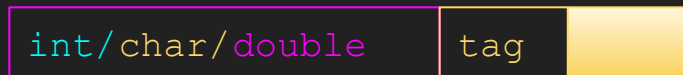
```
tuple<int, char, double>
```



```
union {int i; char c; double d};
```



```
variant<int, char, double>
```



What is it Good for ?

- State Machines
- Value-Semantics for Dynamic Types
 - Commands
- Success/fail
 - `expected<T>`
- Exists/void
 - `optional<T>`
- Run-Time Dispatch (Polymorphism)
- Pattern-Matching

State Machines

```
struct Circling {  
    double mTimeSinceLastShot = ENEMY_SHOT_DELAY;  
    int mNextCirclePosIndex = 0;  
};  
  
struct FlyToCenter { };  
  
struct ShootingFromCenter {  
    double mTimeSinceLastShot = ENEMY_SHOT_DELAY;  
    double mTimeSpentInCenter = 0;  
};  
  
struct FlyOut {  
    int mTargetCornerIndex;  
};  
  
using State = std::variant<  
    Circling,  
    FlyToCenter,  
    ShootingFromCenter,  
    FlyOut>;
```

Meeting C++ 2018
Nikolai Wuttke
std::variant and
the power of
pattern matching

More State Machines

cppcon | 2018

THE C++ CONFERENCE • BELLEVUE, WASHINGTON

EVENTS AND STATES

EVENTS

```
struct event_connect { std::string_view address; };  
struct event_connected {};  
struct event_disconnect {};  
struct event_timeout {};  
using event = std::variant<event_connect, event_connected, event_disconnect, event_timeout>;
```

STATES

```
struct state_idle {};  
struct state_connecting {  
    static constexpr int n_max = 3;  
    int n = 0;  
    std::string address;  
};  
struct state_connected {};  
using state = std::variant<state_idle, state_connecting, state_connected>;
```



MATEUSZ PUSZ

Effective replacement of
dynamic polymorphism
with std::variant

Commands

cppcon | 2016

THE C++ CONFERENCE • BELLEVUE, WASHINGTON



DAVID SANKEL

Variants
Past, Present,
and Future

Command

```
struct command {  
    enum type { SET_SCORE, FIRE_MISSILE, FIRE_LASER, ROTATE };  
    virtual type getType();  
};  
  
struct set_score : command {  
    type getType() { return SET_SCORE; } override final;  
    double value;  
};
```

Now

```
struct set_score {  
    double value;  
};  
using command = variant<set_score, fire_missile, fire_laser, rotate>;
```



12/68

expected

Background Technologies

- `std::variant` (C++17) or `boost::variant`
 - Gives equal importance to all members
- `std::optional` (C++17), `boost::optional`
 - No extra information in the "null" state
- More exotic: the Maybe/Either monads
- Painfully close to what's needed!

cppcon | 2018

THE C++ CONFERENCE • BELLEVUE, WASHINGTON




ANDREI ALEXANDRESCU

Expect the
expected

CppCon.org


Optional

C++ now 2019
MAY 6-10
cppnow.org



Nevin “:-)” Liber

The Many Variants of
std::variant

Video Sponsorship
Provided By: 

optional

- Templated type
 - Closed sum type
 - Holds at most one of the templated type
 - Refinement of `variant`
 - Easier interface
 - Eg: `*o` (dereference) to access object



Run-Time Dispatch

THE ALTERNATIVES - VISITOR

std::variant

```
template <class... Types>  
class variant;
```

std::visit

```
template <class Visitor, class... Variants>  
constexpr auto visit(Visitor&& vis, Variants&&... vars);
```

Lambda

```
[ captures ]<tparams> ( params ) -> ret { body }
```

CPPCON2019 | VIRTUAL TABLES AND ITS ALTERNATIVES | INBAL LEVI

28/46

Cppcon | 2019
The C++ Conference | cppcon.org



Inbal Levi

**Back to Basics:
Virtual Dispatch
and its Alternatives**

Video Sponsorship Provided By:

ansatz

More Run-Time Dispatch



Klaus

Design Evaluation

| | Addition of shapes (OCP) | Addition of operations (OCP) | Separation of Concerns (SRP) | Ease of Use | Performance |
|----------------|--------------------------|------------------------------|------------------------------|-------------|-------------|
| Enum | 1 | 7 | 5 | 6 | 9 |
| OO | 8 | 2 | 2 | 6 | 6 |
| Visitor | 2 | 8 | 8 | 3 | 1 |
| mpark::variant | 3 | 9 | 9 | 9 | 9 |
| | | | | | |
| | | | | | |
| | | | | | |

1 = very bad, ..., 9 = very good

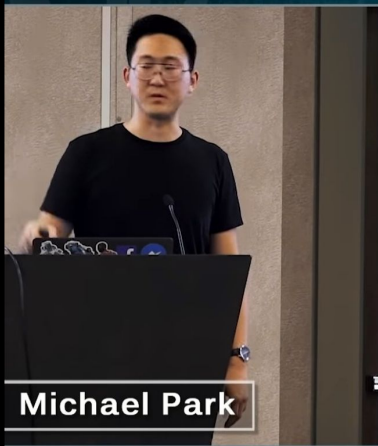
Powered by



StreamYard

90

Pattern Matching



Michael Park

Pattern Matching: A Sneak Peak

Video Sponsorship Provided By:

ansatz

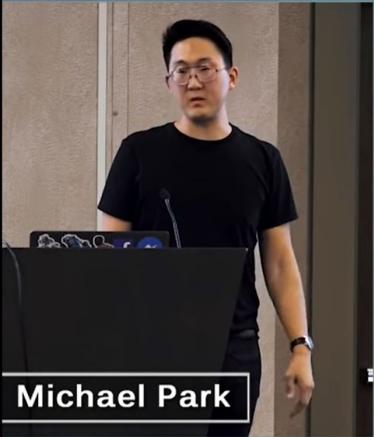
Select + Decompose

```
Message msg = ChangeColor{0, 160, 255};

std::visit(overload{
    [](const Quit&) { std::cout << "Done\n"; },
    [](const Move& move) {
        const auto& [x, y] = move;
        std::cout << "Move by: " << x << ', ' << y << '\n';
    },
    [](const Write& write) {
        const auto& [text] = write;
        std::cout << "Text message: " << text << '\n';
    },
    [](const ChangeColor& change_color) {
        const auto& [r, g, b] = change_color;
        std::cout << "to RGB: " << r << ', ' << g << ', ' << b << '\n';
    }
}, msg);
```

20


Pattern Matching - improved visit()



Michael Park

Pattern Matching:
A Sneak Peak

Video Sponsorship Provided By:



Cppcon | 2019
The C++ Conference | cppcon.org

Bit Bashing [Papers](#) [About the author](#)

std::visit is everything wrong with modern C++


Sep 14, 2017

[HTTPS://BITBASHING.IO/STD-VISIT.HTML](https://bitbashing.io/std-visit.html)

10


Pattern Matching - Anticipated

C++ now 2019
MAY 6-10
cppnow.org





Nevin ":-)" Liber

The Many Variants of
std::variant



Video Sponsorship
Provided By: 

std::variant v1

- Visitation
 - Bjarne Stroustrup
 - Visitation is unpleasant
 - Lots of requests for extension
 - Should just use pattern matching
 - Language feature
 - Not proposed
 - Not for Library Fundamentals TS
 - Shouldn't wait for pattern matching
 - Axel agreed to trim visitation to the bare minimum for variant



Pattern Matching vs. Concepts/Contracts



Sean Parent

Keynote: Better Code: Relationships

Pattern Matching

- Concepts are used as a compile time constraint to select an appropriate operation
- Contracts assert at runtime if an operations preconditions are not met
- A runtime constraint to select an appropriate operation is known as *pattern matching*

```
void f(auto i) requires requires { !(i < 0) }  
void f(int i) [[expects !(i < 0)]]  
void f(int i) requires !(i < 0) // Not yet in C++...
```

Pattern Matching - inspect()

| Before | After |
|---|--|
| <pre>switch (x) { case 0: std::cout << "got zero"; break; case 1: std::cout << "got one"; break; default: std::cout << "don't care"; }</pre> | <pre>inspect (x) { 0 => { std::cout << "got zero"; } 1 => { std::cout << "got one"; } __ => { std::cout << "don't care"; } };</pre> |
| <pre>auto&& [x, y] = p; if (x == 0 && y == 0) { std::cout << "on origin"; } else if (x == 0) { std::cout << "on y-axis"; } else if (y == 0) { std::cout << "on x-axis"; } else { std::cout << x << ', ' << y; }</pre> | <pre>inspect (p) { [0, 0] => { std::cout << "on origin"; } [0, y] => { std::cout << "on y-axis"; } [x, 0] => { std::cout << "on x-axis"; } [x, y] => { std::cout << x << ', ' << y; } };</pre> |

inspect() variants

| Before | After |
|--|---|
| <pre>struct visitor { void operator()(int i) const { os << "got int: " << i; } void operator()(float f) const { os << "got float: " << f; } std::ostream& os; }; std::visit(visitor{strm}, v);</pre> | <pre>inspect (v) { <int> i => { strm << "got int: " << i; } <float> f => { strm << "got float: " << f; } };</pre> |

Inspect() in a nutshell

- Runtime matching types and values
 - `switch()` + `std::visit()` + `[structured, bindings]`
- tuple, variant, any, polymorphic objects, extractors
- First-match, recursive inspection, no fallthrough, `[[strict]]`
- Second tier priority for C++23 (should make progress)

Pattern Matching FTW

```

template <typename T>
void Node<T>::balance() {
    *this = inspect (*this) {
        // left-left case
        //
        //      (Black) z          (Red) y
        //      /      \        /      \
        //    (Red) y    d    (Black) x  (Black) z
        //   /      \    -> /      \    /      \
        // (Red) x    c    a      b    c      d
        // /      \
        // a      b
        [case Black, (*?) [case Red, (*?) [case Red, a, x, b], y, c], z, d]
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
              y,
              std::make_shared<Node>(Black, c, z, d)};
        [case Black, (*?) [case Red, a, x, (*?) [case Red, b, y, c]], z, d] // left-right case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
              y,
              std::make_shared<Node>(Black, c, z, d)};
        [case Black, a, x, (*?) [case Red, (*?) [case Red, b, y, c], z, d]] // right-left case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
              y,
              std::make_shared<Node>(Black, c, z, d)};
        [case Black, a, x, (*?) [case Red, b, y, (*?) [case Red, c, z, d]]] // right-right case
        => Node{Red, std::make_shared<Node>(Black, a, x, b),
              y,
              std::make_shared<Node>(Black, c, z, d)};
        self => self; // do nothing
    };
}

```

Variants vs. Unions

Tag is private → variant is safe

- Only constructor and assignment-operator change the tag.
- No chance of user error
- Compiler knows all the alternatives

```
union IdentityCard {  
    IDNumber nationalID;  
    PassportNumber passport;  
    UUID factoryCertificate;  
};  
  
enum IDType { CITIZEN, TOURIST, ROBOT };  
void checkID(IdentityCard id, IDType type) {  
    switch (type) {  
        case TOURIST: checkPassport(id.passport);  
        case CITIZEN: checkPassport(id.passport);  
    }  
}
```

ilbarkan@gmail.com

Bugs in this C Code:



Access the incorrect member



Forget missing types



Forget to break

```
using IdentityCard = variant<IDNumber,  
                             PassportNumber, UUID>;  
  
void checkID(IdentityCard id) {  
    visit([](auto& x) {x.check();}, id);  
}
```


How do they Actually do it in C

tag

int/char/double

- Explicit Header:

```
struct IdentityCard {
    IDType type;
    union /*value*/ {
        IDNumber id;
        PassportNumber passport;
    };
};
```

- Implicit Header:

```
union IdentityCard {
    struct Header { IDType type; };
    struct Citizen {IDType type;
        IDNumber id;
    };
    struct Tourist {IDType type;
        PassportNumber passport;
    };
};
```

- Over time header gets more data
 - Expiration date, Photo, ...
- Type specific code need access to header.
- Preprocessor (??):

```
#define HEADER_FIELDS \
    IDType type; \
    Date expiration;

struct Header {  HEADER_FIELDS };
struct Citizen {  HEADER_FIELDS
    IDNumber id;
};
struct Tourist { HEADER_FIELDS
    PassportNumber id;
};
union IdentityCard {
    Header header;
    Citizen citizen;
    Tourist tourist;
};
```

Keeping the C Layout

tag

int/char/double

- Header with type is common
 - Network Protocols - TCP/IP, Finance
 - File Formats - ELF
 - Serialization - Cap'n Proto, Apache Avro
- C layout is important
 - Compatibility with existing code
- Goal - Be safer than C, keep the layout
 - Sacrifice some safety

Unions in the C++ Standard

tag

int/char/double

- Accessing header fields is properly defined in many cases.
- Standard-layout class is part of the C++ standard
 - “Standard-layout classes are useful for communicating with code written in other programming languages.”
 - Various constraints (no virtuals, single access control, all non-statics in same class,...)
 - C++11: `static_assert(std::is_standard_layout<Citizen>::value, "not standard");`
 - C++17: `static_assert(std::is_standard_layout_v<Citizen>);`
- Layout-compatibility allows accessing members without knowing the type
 - “In a standard-layout union with an active member of struct type T1, it is permitted to read a non-static data member m of another union member of struct type T2 provided m is part of the common initial sequence of T1 and T2”
 - C++20: `static_assert(std::is_corresponding_member(&Header::type, &Citizen::type));`

intrusive_variant

tag

int/char/double

```
using ID_intrusive = intrusive_variant<
    IDType, offsetof(IDHeader, type),
    intrusive_variant_tag_type<IDType::CITIZEN,Citizen>,
    intrusive_variant_tag_type<IDType::TOURIST,Tourist>>;
// ... Alternatively ...
using ID_intrusive = intrusive_variant<
    IDType, offsetof(IDHeader, type),
    intrusive_variant_type<Citizen>,
    intrusive_variant_type<Citizen>>;
```

- User dictates the type and location of the tag
- visit() is still $O(1)$
 - Potentially larger lookup table
- Customization Point for tag deduction

```
union IdentityCard {
    struct IDHeader { {
        const IDType type;
    } header;
    struct Citizen {
        const IDType type;
        static constexpr IDType TAG =
            IDType::CITIZEN;
        IDNumber id;
    } citizen;t;
    struct Tourist {
        const IDType type;
        static constexpr IDType TAG =
            IDType::TOURIST;
        PassportNumber passport;
    } tourist;
};
```

Different Approaches to get the Tag

- Offset of the field in the object

```
getTag<IDType>(hdr, offsetof(Hdr, tag));  
getTag<IDType>(hdr, std::integral_constant<size_t, offsetof(Hdr, tag)>());
```

- Pointer to the field

```
getTag<IDType>(hdr, &Hdr::tag);
```

- Call a member function

```
getTag<IDType>(hdr, &Hdr::getTag);  
    ○ Useful when the tag is private  
    ○ Useful when the tag is calculated
```

- Call a free-function / lambda

```
getTag<IDType>(hdr, [](const Hdr& h) { return h.tag; });  
    ○ Less intrusive
```

Implementation of Tag Extraction

```
template <typename R, typename T, typename Offset>
R getTag(const T& x,
        Offset offset) requires std::convertible_to<Offset, size_t> {
    return *reinterpret_cast<const R*>(reinterpret_cast<const char*>(&x) +
                                       offset);
}

template <typename R, typename T, typename M>
R getTag(const T& x, M method) requires std::invocable<M, T> {
    return std::invoke(method, x);
}
```

- `std::invoke()` is flexible - functions, lambdas, pointer-to-member

Add C++20 Safety

tag

int/char/double

```
using ID_intrusive = decltype(decl_safe_intrusive_variant(  
    &Citizen::type, IDType::CITIZEN,  
    &Tourist::type, IDType::TOURIST));
```

- Helper function to ease type deduction.
- Validate corresponding members.
 - No need to deal with offsets, etc.
- constexpr generic lambdas can be used for extra beauty

Questions, Comments...

More C++ Safety

- Intrusive_variant has safe visit() and links the Type with the Tag
- Still - we need to add boilerplate and can still have bugs.
- Class Hierarchies can do better:
 - Base class is essentially a header.
 - Use a (constexpr) lambda to get the tag of each type
- Utilities:
 - is_base_of<> - to make sure all types have the right base
 - decltype() to get to type and static members
- The problem - base class with data members isn't *standard-layout*
 - But it generally works. Perhaps we can update the standard

Sean Parent on variant member similarity



variant_of_base

```
struct IDHeader {
    IDHeader(IDType type) : m_type(type) {}
    const IDType m_type;
};

struct Citizen : public IDHeader {
    Citizen() : IDHeader(TAG) {}
    static constexpr IDType TAG = CITIZEN;
    IDNumber id;
};

struct Tourist : public IDHeader {
    Tourist() : IDHeader(TAG) {}
    static constexpr IDType TAG = TOURIST;
    PassportNumber passport;
};
```

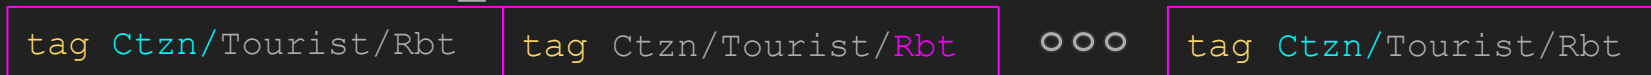
```
using ID_with_base =
    variant_of_base<IDHeader,
                    [] (const IDHeader &hdr) {
                        return hdr.m_type; },
                    [] (auto *x) {
                        return decltype(*x)::TAG;
                    },
                    Citizen, Tourist>;
```

Arrays of Variants

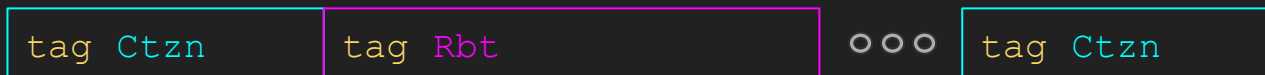
- `vector<variant<Citizen, Tourist, Robot>>`



- `vector<intrusive_variant<Citizen, Tourist, Robot>>`



- C-Style (real world) - use only the RAM we need



condensed_variant utilities

- condensed_variant_iterator: const forward iterator over variants
- condensed_variant_queue: emplace_back/pop_front container
- Root of the logic is knowing how much to jump:

```
const Base *next =  
    visit([](const auto *p) { return static_cast<const Base *>(p + 1); },  
          myVariant);
```

- Key point: (p + 1) knows the correct type.

Summary so far

- variants are different than unions
- real-world unions already have tags (and headers)
- Intrusive_variant - C++ safety with high C compatibility
- Variant_of_base - add classes to your code
 - Not standard-layout, undefined behavior
 - Perhaps we should widen the rules, add `[[standard_layout]]` ?
- condensed_variant - real world streams of binary data

Questions, Comments...

Assisting Devirtualization

Blast from the Past

Virtual dispatch analysis

- + Promotes flexibility & decoupling
- + Best for large, unbounded hierarchies
- + Automatic 'load balancing' of icache
- Wastes space in the hot zone
- Relatively costly
- Performs poorly on small/closed hierarchies
- Can't change object type in-situ
- Pay for *potential*, not *realized* flexibility

Devirtualization, take 2

```
class Base {  
    struct VTable {  
        int (*get)(const Base&);  
        int (*set)(Base&, int);  
    };  
    static VTable vtbl[totalClasses];  
    uint8_t tag;  
public:  
    int get() const {  
        return (vtbl[tag].get)(*this);  
    }  
    int set(int x) {  
        (vtbl[tag].set)(*this, x);  
    }  
};
```

© 2013- Andrei Alexandrescu. Do not redistribute.



The Cost of Virtual Functions

- Everyone is afraid of **Branch Mispredictions**
- However - Processors have relatively good predictors.
 - Processors learn your program **and the data** as it runs..
- Compilers typically only see the program (or part of it)
 - Virtues of PGO, Virtues of LTO
 - C++20: `[[likely]]`
- Devirtualization lets compilers break through virtual calls: inline, inspect, rearrange code.

std::visit for devirtualization

- Base class with some Implementation

```
struct Base {  
    virtual int foo(){};  
};  
struct D1 final : public Base {  
    int foo() override { return 1; };  
};  
struct D2 final : public Base {  
    int foo() override { return 2; };  
};
```

- Variant of pointers to the same hierarchy.

```
variant<D1*,Base*> myVariant;  
visit([](auto* p) {p->foo();}, myVariant);
```

Compiler Explorer



Add...

More

Sponsors

PC-lint

No bugs

Solid Sands

Share

Other

C++ source #1 X

A Save/Load + Add new... Vim CppInsights Quick-bench C++

```
1 #include <variant>
2
3 using std::variant;
4
5 struct Base {
6     virtual ~Base() = default;
7     virtual int foo() = 0;
8 };
9
10 struct Common final : public Base {
11     int foo() override {
12         return 750;
13     }
14 };
15
16 struct Rare final : public Base {
17     int foo() override {return 322;}
18 };
19
20 int test(std::variant<Common*,Base*> pBase) {
21     return visit([](auto*p) { return p->foo();}, pBase);
22 }
```

x86-64 clang (trunk) (Editor #1, Compiler #1) C++ X

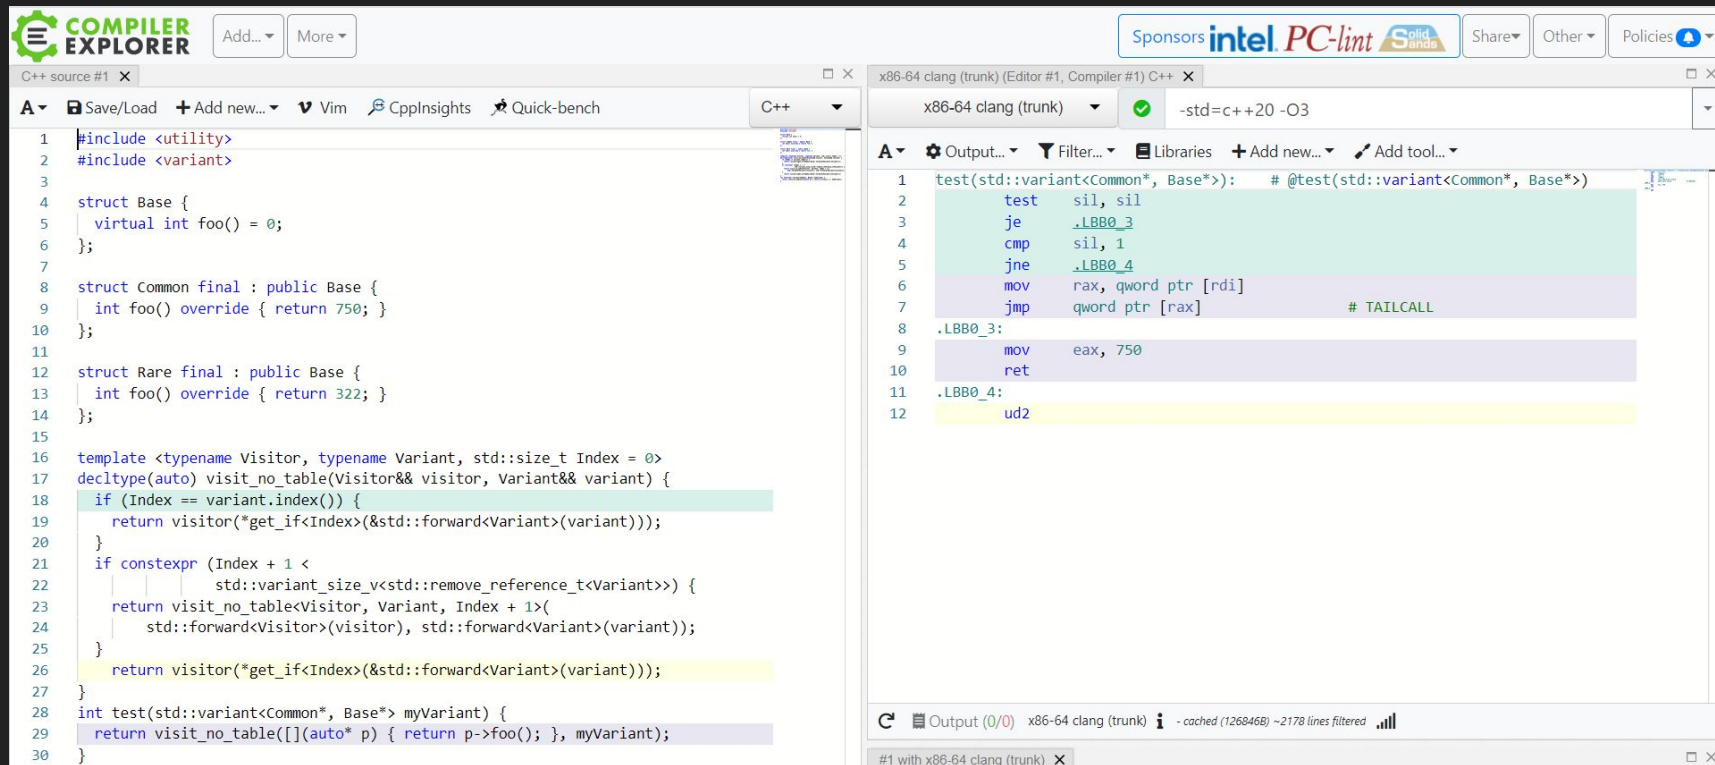
x86-64 clang (trunk)

-O3 -std=c++20

A Output... Filter... Libraries + Add new... Add tool...

```
1 test(std::variant<Common*, Base*>): # @test(std::variant<Common*, Base*>)
2     sub     rsp, 24
3     mov     qword ptr [rsp + 8], rdi
4     mov     byte ptr [rsp + 16], sil
5     cmp     sil, -1
6     movzx   eax, sil
7     mov     rcx, -1
8     cmovne  rcx, rax
9     mov     rdi, rsp
10    lea     rsi, [rsp + 8]
11    call    qword ptr [8*rcx + std::__detail::__variant::__gen_vtable<tru
12    add     rsp, 24
13    ret
14 std::__detail::__variant::__gen_vtable_impl<true, std::__detail::__variant::
15     mov     eax, 750
16     ret
17 std::__detail::__variant::__gen_vtable_impl<true, std::__detail::__variant::
18     mov     rdi, qword ptr [rsi]
19     mov     rax, qword ptr [rdi]
20     jmp     qword ptr [rax + 16] # TAILCALL
21 std::__detail::__variant::__gen_vtable<true, int, test(std::variant<Common*,
22     .quad   std::__detail::__variant::__gen_vtable_impl<true, std::__deta
23     .quad   std::__detail::__variant::__gen_vtable_impl<true, std::__deta
```

“Slower” visit (<https://godbolt.org/z/eTcEa5oqE>)



The screenshot displays the Compiler Explorer interface. The left pane shows the C++ source code for a visitor pattern. The right pane shows the assembly output generated by x86-64 clang (trunk) with flags `-std=c++20 -O3`.

C++ source code (C++ source #1):

```
1 #include <utility>
2 #include <variant>
3
4 struct Base {
5     virtual int foo() = 0;
6 };
7
8 struct Common final : public Base {
9     int foo() override { return 750; }
10 };
11
12 struct Rare final : public Base {
13     int foo() override { return 322; }
14 };
15
16 template <typename Visitor, typename Variant, std::size_t Index = 0>
17 decltype(auto) visit_no_table(Visitor&& visitor, Variant&& variant) {
18     if (Index == variant.index()) {
19         return visitor(*get_if<Index>>(&std::forward<Variant>(variant)));
20     }
21     if constexpr (Index + 1 <
22                   std::variant_size_v<std::remove_reference_t<Variant>>) {
23         return visit_no_table<Visitor, Variant, Index + 1>(
24             std::forward<Visitor>(visitor), std::forward<Variant>(variant));
25     }
26     return visitor(*get_if<Index>>(&std::forward<Variant>(variant)));
27 }
28 int test(std::variant<Common*, Base*> myVariant) {
29     return visit_no_table([](auto* p) { return p->foo(); }, myVariant);
30 }
```

Assembly output (x86-64 clang (trunk) (Editor #1, Compiler #1) C++):

```
1 test(std::variant<Common*, Base*>): # @test(std::variant<Common*, Base*>)
2     test    sil, sil
3     je      .LBB0_3
4     cmp     sil, 1
5     jne     .LBB0_4
6     mov     rax, qword ptr [rdi]
7     jmp     qword ptr [rax] # TAILCALL
8 .LBB0_3:
9     mov     eax, 750
10    ret
11 .LBB0_4:
12    ud2
```

The assembly output shows a tail call to the `visit_no_table` function for the `Rare` variant, which is not shown in the snippet. The `Common` variant returns 750.

Thank You